


```
[robert.campbell@apex wk6]$ cat pi_mc_mpi.out
```

npoints	pi	nprocs	elapsed wall-clock time
10000000	3.1418	1	1.15979
10000000	3.1414	2	0.634331
10000000	3.14148	4	0.348533
10000000	3.14205	8	0.176174
10000000	3.1419	16	0.0885261
10000000	3.14146	32	0.105811

```
[robert.campbell@apex wk6]$ cat pi_mc_mpi.out
```

npoints	pi	nprocs	elapsed wall-clock time
100000000	3.14156	1	5.88974
100000000	3.14137	2	3.57555
100000000	3.14158	4	2.01964
100000000	3.14151	8	1.24406
100000000	3.14176	16	0.857434
100000000	3.14177	32	0.571592



Both data sets saw massive improvements with diminishing returns as additional threads were added, with  $n=100,000,000$  having benefitting more for each additional thread. One thing to notice is that the smaller sample size was less precise in its result, with the 3<sup>rd</sup> decimal place fluctuating from the correct value while the larger sample size fluctuated around the 4<sup>th</sup> decimal place. As far as this lab goes, it looks like 1 additional order of magnitude of samples added an order of magnitude of precision in the result, however it did noticeably increase serial runtime (~6x). If a more precise answer is needed, even more samples would be required, which would benefit even more from parallelization.