

# Dataframe basics

Bryan Clair

Wednesday, February 3, 2021

# Statistical Experiments

Say something true about the world, even though we only observe a sample.

## Key idea

- Introduce randomness into our observations and use **probability** to make precise statements.

A **statistical experiment** is a repeatable process that produces a random outcome.

The **sample space** is the set of all possible outcomes.

# Sample Space

Sample space: the set of outcomes of a statistical experiment.

Flip a coin

$S = \text{Heads, Tails}$

Roll a die

$S = 1, 2, 3, 4, 5, 6$

Wait for a bus at a bus stop

$S = [0, \infty)$

Take a patient's blood pressure

$S$  is pairs of positive numbers

# Events

An **event** is a subset of the sample space.

Flip a coin. There are four possible events:


$\{\}$ ,  $\{\text{Heads}\}$ ,  $\{\text{Tails}\}$ ,  $\{\text{Heads}, \text{Tails}\}$

Roll a die. The event “even” is 2, 4, 6.

Wait for a bus. The event “wait less than five minutes” is  $[0, 5)$ .

# Events

Check blood pressure. The event “normal blood pressure” is:



# Probability

The **probability of an event** is a number between 0 and 1.

For an event  $A$ , write  $P(A)$  for the probability.

$$0 \leq P(A) \leq 1$$

We *assign* probabilities to events in the hopes that they model the world.

The Law of Large Numbers:

$$\frac{\text{\# of times } A \text{ occurs}}{\text{\# of trials}} \rightarrow P(A)$$

as the number of trials of the experiment grows.

# Simulation

We will use R to estimate probabilities by simulation.

**sample**: choose randomly from a vector

**replicate**: repeat an experiment

# Examples

Roll a die. What is  $P(\text{roll} < 3)$ ?

Roll two dice and sum them.

- What is  $P(\text{roll} = 7)$ ?
- What is  $P(\text{roll} = 12)$ ?
- What is  $P(\text{roll is even})$ ?

Roll three dice and sum them. What is  $P(\text{roll} > 12)$ ?

Flip 20 coins. What is the probability of getting exactly 10 heads?

Sample the lengths of four rivers from the `rivers` vector. What is the probability that your sample average is larger than 1000?



# More Examples

10% of humans are left handed. Do we have an unusual number of lefties in this class?

If a 75% free throw shooter takes 10 free throws, what is the probability they make them all?

What is the probability they make more than half?